

# R Scripting

## Exercises for unit 2 - Data structures (II)

Marcus Wurzer

Please solve the following problems!

1. Try to answer quiz questions 4. and 5. at the beginning of the “Data structures” chapter of the coursebook: <http://adv-r.had.co.nz/Data-structures.html>
2. Try to answer quiz questions 1. to 5. at the beginning of the “Subsetting” chapter of the coursebook: <http://adv-r.had.co.nz/Subsetting.html>
3. Coursebook exercises:

- a. What are the six types of atomic vector? How does a list differ from an atomic vector?
- b. Test your knowledge of vector coercion rules by predicting the output of the following uses of `c()`:

```
c(1, FALSE)
c("a", 1)
c(list(1), "a")
c(TRUE, 1L)
```

- c. Why is the default missing value, `NA`, a logical vector? What’s special about logical vectors? (Hint: think about `c(FALSE, NA_character_)`.)
- d. What does `dim()` return when applied to a vector?
- e. How would you describe the following three objects? What makes them different to `1:5`?

```
x1 <- array(1:5, c(1, 1, 5))
x2 <- array(1:5, c(1, 5, 1))
x3 <- array(1:5, c(5, 1, 1))
```

- f. What does `as.matrix()` do when applied to a data frame with columns of different types?
- g. Fix each of the following common data frame subsetting errors:

```
mtcars[mtcars$cyl = 4, ]
mtcars[-1:4, ]
mtcars[mtcars$cyl <= 5]
mtcars[mtcars$cyl == 4 | 6, ]
```

- h. Why does `mtcars[1:20]` return an error? How does it differ from the similar `mtcars[1:20, ]`?
- i. What does `df[is.na(df)] <- 0` do? How does it work?
- j. How would you randomly permute the columns of a data frame? (This is an important technique in random forests.) Can you simultaneously permute the rows and columns in one step?
- k. How could you put the columns in a data frame in alphabetical order?

4. The following table is given:

	name	figure	height	weight	drinks	job
1	Susi	chubby	1.97	98	TRUE	student
2	Tim	chubby	1.67	89	n/a	student
3	Christine	beefy	1.90	71	TRUE	student
4	Mathias	skinny	1.81	86	TRUE	employed

(n/a = not available)

- a. Generate the variables

**name** as character string

**job** as factor with the following categories: *student, employed, self-employed*

**figure** as ordered factor (using function `ordered()`) with the following categories: *skinny, lean, slim, normal, chubby, beefy*

**height** as numeric (continuous) variable

**weight** as numeric (discrete) variable

**drinks** as logical (binary) variable

Pay attention to the correct specification of missing values!

- b. Generate a data frame **friends** from these variables that looks similar to the table given above and print it!

5. Dataset **chickwts** is included in R:

```
data(chickwts)
?chickwts
```

It is a data frame that contains two variables: The weight of chickens contingent upon various feed supplements (**weight** and **feed**).

- a. Print the first 6 rows of the dataset.
- b. Extract all factor levels of the variable **feed**.
- c. Extract all rows of the data frame for chickens that are fed with **meatmeal** (Hint: `subset()`).
- d. Extract the weight of the chickens that are fed with **casein** or **horsebean** (Hint: Operator for logical or: `|`).
- e. Compute the mean chicken weight separately for each feed supplement (Hint: `aggregate()`).

6. Use dataset **chickwts** again.

- a. Compute the mean weight of all chickens (Hint: function `mean()`).
- b. Use **weight** to generate a categorical variable **weight\_cat** with the three categories **light**, **medium**, and **heavy**. All categories should approximately contain the same number of chickens (Hint: `cut()`).
- c. Subsequently, combine the **light** and **medium** categories to a new category **standard**, and rename the **heavy** category to **premium**.
- d. How many “premium” chickens are there that have been fed with **meatmeal**? (Do not count manually, but use R functions!)

7. A digitized black and white image can be represented by a binary matrix, with “0” symbolizing color value “white” and “1” color value “black”.

- a. Create an 800x600 matrix and randomly fill it with black and white values.
- b. Compute the grey value (= mean color value) of the whole matrix.
- c. Compute the grey value for a certain image section, namely the lower right quadrant of the picture (Hint: Use horizontal and vertical splitting to get the four equally sized image sections).
- d. Invert the picture to get a negative (0 becomes 1 and vice versa). Compute the grey value again.

8. Use the following R-code to generate a matrix of temperature measurements (in °C) for the 31 daily average high temperatures of July in two cities:

```

set.seed(1)
temps <- matrix(c(round(rnorm(31, 26, 5), 1), round(rnorm(31, 22, 4), 1)),
                 nrow = 2, byrow = TRUE)
rownames(temps) <- c("City A", "City B")
colnames(temps) <- paste("July", 1:31)

```

Unfortunately, many of the entries have been lost when reading the data...:

```

temps[sample(1:length(temps), 20)] <- NA

```

- a. Compute mean, minimum, and maximum for all entries and separately for both cities.
- b. Compute the maximum temperature for the second half of the month (starting with the 16<sup>th</sup> of July)
- c. Save the observational units in a list having two components (one for each city). The list should only contain the actual measurements and reference dates (i.e., without missings).