

R Scripting

Exercises for unit 2 - Data structures (II)

Marcus Wurzer

Please solve the following problems!

1. Try to answer quiz questions 4. and 5. at the beginning of the “Data structures” chapter of the coursebook: <http://adv-r.had.co.nz/Data-structures.html>

See coursebook for answers!

2. Try to answer quiz questions 1. to 5. at the beginning of the “Subsetting” chapter of the coursebook: <http://adv-r.had.co.nz/Subsetting.html>

See coursebook for answers!

3. Coursebook exercises:

- a. What are the six types of atomic vector? How does a list differ from an atomic vector?

The six types are logical, integer, double, character, complex and raw. The elements of a list don't have to be of the same type.

- b. Test your knowledge of vector coercion rules by predicting the output of the following uses of `c()`:

```
c(1, FALSE)      # will be coerced to numeric    -> 1 0
c("a", 1)        # will be coerced to character -> "a" "1"
c(list(1), "a")   # will be coerced to a list with two elements of type double and character
c(TRUE, 1L)      # will be coerced to integer   -> 1 1
```

- c. Why is the default missing value, `NA`, a logical vector? What's special about logical vectors? (Hint: think about `c(FALSE, NA_character_)`.)

It is a practical thought. When you combine `NA`s in `c()` with other atomic types they will be coerced like `TRUE` and `FALSE` to integer (`NA_integer_`), double (`NA_real_`), complex (`NA_complex_`) and character (`NA_character_`). Recall that in R there is a hierarchy of coercion that goes logical -> integer -> double -> character. If `NA` were, for example, a character, including `NA` in a set of integers or logicals would result in them getting coerced to characters which would be undesirable. Making `NA` a logical means that involving an `NA` in a dataset (which happens often) will not result in coercion.

- d. What does `dim()` return when applied to a vector?

`NULL`

- e. How would you describe the following three objects? What makes them different to `1:5`?

```
x1 <- array(1:5, c(1, 1, 5)) # 1 row, 1 column, 5 in third dimension
x2 <- array(1:5, c(1, 5, 1)) # 1 row, 5 columns, 1 in third dimension
x3 <- array(1:5, c(5, 1, 1)) # 5 rows, 1 column, 1 in third dimension
```

They are of class `array` and so they have a `dim` attribute.

- f. What does `as.matrix()` do when applied to a data frame with columns of different types?

From `?as.matrix`: The method for data frames will return a character matrix if there is only atomic columns and any non-(numeric/logical/complex) column, applying `as.vector` to factors and `format` to other non-character columns. Otherwise the usual coercion hierarchy (logical < integer < double < complex) will be used, e.g., all-logical data frames will be coerced to a logical matrix, mixed logical-integer will give a integer matrix, etc.

- g. Fix each of the following common data frame subsetting errors:

```
mtcars[mtcars$cyl = 4, ]      # = -> ==
mtcars[-1:4, ]               # -1:4 -> -(1:4)
mtcars[mtcars$cyl <= 5]      # ", " is missing
mtcars[mtcars$cyl == 4 | 6, ] # 6 -> mtcars$cyl == 6
```

- h. Why does `mtcars[1:20]` return an error? How does it differ from the similar `mtcars[1:20,]`?

In the first case `mtcars` is subsetted with a vector and the statement should return a `data.frame` of the first 20 columns in `mtcars`. Since `mtcars` only has 11 columns, the index is out of bounds, which explains the error. The biggest difference of `mtcars[1:20,]` to the former case, is that now `mtcars` is subsetted with two vectors. In this case you will get returned the first 20 rows and all columns (like subsetting a matrix).

- i. What does `df[is.na(df)] <- 0` do? How does it work?

It replaces all NAs within `df` with the value 0. `is.na(df)` returns a logical matrix which is used to subset `df`. Since you can combine subsetting and assignment, only the matched part of `df` (the NAs) is replaced with 0 entries.

- j. How would you randomly permute the columns of a data frame? (This is an important technique in random forests.) Can you simultaneously permute the rows and columns in one step?

```
iris[sample(ncol(iris))] # permute rows
iris[sample(nrow(iris)), sample(ncol(iris)), drop = FALSE] # permute both at the same time
```

- k. How could you put the columns in a data frame in alphabetical order?

```
iris[sort(names(iris))]
```

4. The following table is given:

	name	figure	height	weight	drinks	job
1	Susi	chubby	1.97	98	TRUE	student
2	Tim	chubby	1.67	89	n/a	student
3	Christine	beefy	1.90	71	TRUE	student
4	Mathias	skinny	1.81	86	TRUE	employed

(n/a = not available)

- a. Generate the variables

name as character string

job as factor with the following categories: *student*, *employed*, *self-employed*

figure as ordered factor (using function `ordered()`) with the following categories: *skinny*, *lean*, *slim*, *normal*, *chubby*, *beefy*

height as numeric (continuous) variable

weight as numeric (discrete) variable

drinks as logical (binary) variable

Pay attention to the correct specification of missing values!

- b. Generate a data frame `friends` from these variables that looks similar to the table given above and print it!

```
# a.
name <- c("Susi", "Tim", "Christine", "Mathias")
job <- factor(c(rep("student", 3), "employed"),
              levels = c("student", "employed", "self-employed"))
figure <- ordered(c("chubby", "chubby", "beefy", "skinny"),
                  levels = c("skinny", "lean", "slim", "normal", "chubby", "beefy"))
height <- c(1.97, 1.67, 1.90, 1.81)
weight <- c(98L, 89L, 71L, 86L)
drinks <- c(TRUE, NA, TRUE, TRUE)

# b.
friends <- data.frame(name, figure, height, weight, drinks, job)
friends
```

```
##      name figure height weight drinks  job
## 1    Susi chubby  1.97    98    TRUE student
## 2     Tim chubby  1.67    89     NA student
## 3 Christine beefy  1.90    71    TRUE student
## 4  Mathias skinny  1.81    86    TRUE employed
```

5. Dataset `chickwts` is included in R:

```
data(chickwts)
?chickwts
```

It is a data frame that contains two variables: The weight of chickens contingent upon various feed supplements (weight and feed).

- Print the first 6 rows of the dataset.
- Extract all factor levels of the variable `feed`.
- Extract all rows of the data frame for chickens that are fed with `meatmeal` (Hint: `subset()`).
- Extract the weight of the chickens that are fed with `casein` or `horsebean` (Hint: Operator for logical or: `|`).
- Compute the mean chicken weight separately for each feed supplement (Hint: `aggregate()`).

```
# a.
head(chickwts)

##   weight      feed
## 1   179 horsebean
## 2   160 horsebean
## 3   136 horsebean
## 4   227 horsebean
## 5   217 horsebean
## 6   168 horsebean

# b.
levels(chickwts$feed)

## [1] "casein"      "horsebean" "linseed"   "meatmeal"  "soybean"   "sunflower"

# c.
subset(chickwts, feed == "meatmeal")

##   weight      feed
## 49   325 meatmeal
## 50   257 meatmeal
```

```
## 51    303 meatmeal
## 52    315 meatmeal
## 53    380 meatmeal
## 54    153 meatmeal
## 55    263 meatmeal
## 56    242 meatmeal
## 57    206 meatmeal
## 58    344 meatmeal
## 59    258 meatmeal
```

```
# d.
subset(chickwts, feed == "casein" | feed == "horsebean", weight)
```

```
##      weight
## 1         179
## 2         160
## 3         136
## 4         227
## 5         217
## 6         168
## 7         108
## 8         124
## 9         143
## 10        140
## 60        368
## 61        390
## 62        379
## 63        260
## 64        404
## 65        318
## 66        352
## 67        359
## 68        216
## 69        222
## 70        283
## 71        332
```

```
# or
subset(chickwts, feed %in% c("casein", "horsebean"), weight)
```

```
##      weight
## 1         179
## 2         160
## 3         136
## 4         227
## 5         217
## 6         168
## 7         108
## 8         124
## 9         143
## 10        140
## 60        368
## 61        390
## 62        379
## 63        260
## 64        404
```

```
## 65    318
## 66    352
## 67    359
## 68    216
## 69    222
## 70    283
## 71    332
```

```
# e.
aggregate(weight ~ feed, chickwts, mean)
```

```
##      feed  weight
## 1  casein 323.5833
## 2 horsebean 160.2000
## 3  linseed 218.7500
## 4  meatmeal 276.9091
## 5   soybean 246.4286
## 6 sunflower 328.9167
```

6. Use dataset `chickwts` again.

- Compute the mean weight of all chickens (Hint: function `mean()`).
- Use `weight` to generate a categorical variable `weight_cat` with the three categories `light`, `medium`, and `heavy`. All categories should approximately contain the same number of chickens (Hint: `cut()`.)
- Subsequently, combine the `light` and `medium` categories to a new category `standard`, and rename the `heavy` category to `premium`.
- How many “premium” chickens are there that have been fed with `meatmeal`? (Do not count manually, but use R functions!)

```
# a.
mean(chickwts$weight)
```

```
## [1] 261.3099
```

```
# b.
nrow(chickwts) # approx. 24 should be in each group
```

```
## [1] 71
```

```
idx <- round(nrow(chickwts) * c(1/3, 2/3)) # divide n by 3 to get
# the indices of the breaks
breaks <- sort(chickwts$weight)[idx] # sort the chickens by weight
# and get the breaks
weight_cat <- cut(chickwts$weight, c(-Inf, breaks, Inf),
                  labels = c("light", "medium", "heavy"))
table(weight_cat) # show frequencies of the categories
```

```
## weight_cat
## light medium heavy
##    24    23    24
```

```
# c.
levels(weight_cat)[1:2]
```

```
## [1] "light" "medium"
```

```
levels(weight_cat)[1:2] <- "standard"
levels(weight_cat)[2] <- "premium"
```

```
table(weight_cat)

## weight_cat
## standard premium
##      47      24

chickwts.prem <- chickwts[weight_cat == "premium", ]
summary(chickwts.prem$feed) # -> 4 chickens

## casein horsebean linseed meatmeal soybean sunflower
##      8      0      0      4      3      9

# or
chickwts.prem <- subset(chickwts,
                        weight_cat == "premium" & feed == "meatmeal")
nrow(chickwts.prem)

## [1] 4
```

7. A digitized black and white image can be represented by a binary matrix, with “0” symbolizing color value “white” and “1” color value “black”.
- Create an 800x600 matrix and randomly fill it with black and white values.
 - Compute the grey value (= mean color value) of the whole matrix.
 - Compute the grey value for a certain image section, namely the lower right quadrant of the picture (Hint: Use horizontal and vertical splitting to get the four equally sized image sections).
 - Invert the picture to get a negative (0 becomes 1 and vice versa). Compute the grey value again.

```
# a.
X <- matrix(sample(0:1, 800 * 600, replace = TRUE), nrow = 800)
dim(X)

## [1] 800 600

table(X)

## X
##      0      1
## 239273 240727

# b.
mean(X)

## [1] 0.5015146

# c.
row.idx <- (nrow(X) / 2 + 1):nrow(X)
col.idx <- (ncol(X) / 2 + 1):ncol(X)
X_lr <- X[row.idx, col.idx]
mean(X_lr)

## [1] 0.5017833

# d.
X_i <- -1 * X + 1
mean(X_i)

## [1] 0.4984854
```

8. Use the following R-code to generate a matrix of temperature measurements (in °C) for the 31 daily average high temperatures of July in two cities:

```

set.seed(1)
temps <- matrix(c(round(rnorm(31, 26, 5), 1), round(rnorm(31, 22, 4), 1)),
                nrow = 2, byrow = TRUE)
rownames(temps) <- c("City A", "City B")
colnames(temps) <- paste("July", 1:31)

```

Unfortunately, many of the entries have been lost when reading the data...:

```

temps[sample(1:length(temps), 20)] <- NA

```

- Compute mean, minimum, and maximum for all entries and separately for both cities.
- Compute the maximum temperature for the second half of the month (starting with the 16th of July)
- Save the observational units in a list having two components (one for each city). The list should only contain the actual measurements and reference dates (i.e., without missings).

```

# a.
mean(temps, na.rm = TRUE)

## [1] 24.5881

min(temps, na.rm = TRUE)

## [1] 14.9

max(temps, na.rm = TRUE)

## [1] 34

mean(temps[1, ], na.rm = TRUE)

## [1] 26.26667

mean(temps[2, ], na.rm = TRUE)

## [1] 22.35

min(temps[1, ], na.rm = TRUE)

## [1] 14.9

min(temps[2, ], na.rm = TRUE)

## [1] 16.5

max(temps[1, ], na.rm = TRUE)

## [1] 34

max(temps[2, ], na.rm = TRUE)

## [1] 29.9

# b.
max(temps[, 16:ncol(temps)], na.rm = TRUE)

## [1] 32.8

# c.
temps.list <- list("City A" = na.omit(temps[1, ]),
                  "City B" = na.omit(temps[2, ]))
temps.list

```

```

## $'City A'
##   July 2   July 3   July 4   July 5   July 6   July 7   July 9   July 10   July 11   July 12
##   26.9    21.8    34.0    27.6    21.9    28.4    28.9    24.5    33.6    27.9
##   July 13   July 14   July 16   July 18   July 19   July 21   July 23   July 24   July 25   July 27
##   22.9    14.9    25.8    30.7    30.1    30.6    26.4    16.1    29.1    25.2
##   July 28   July 29   July 30   July 31
##   18.6    23.6    28.1    32.8
## attr(,"na.action")
##   July 1   July 8   July 15   July 17   July 20   July 22   July 26
##     1       8      15       17      20      22      26
## attr(,"class")
## [1] "omit"
##
## $'City B'
##   July 1   July 2   July 3   July 4   July 7   July 8   July 9   July 10   July 15   July 16
##   21.6    23.6    21.8    16.5    21.8    26.4    25.1    21.3    19.2    23.5
##   July 18   July 19   July 22   July 23   July 25   July 27   July 28   July 29
##   21.6    25.5    23.4    17.5    29.9    17.8    24.3    21.5
## attr(,"na.action")
##   July 5   July 6   July 11   July 12   July 13   July 14   July 17   July 20   July 21   July 24
##     5       6      11      12      13      14      17      20      21      24
##   July 26   July 30   July 31
##     26      30      31
## attr(,"class")
## [1] "omit"

```